

5.1 Common Pitfalls: Why Mood Keywords and Stage Directions Fail

Before building the filter, confirm you are not defaulting to one of these three structurally broken approaches. Each one fails for a specific reason that the Three-Tier Emotional Filter is designed to solve.

Approach	What You Did	Why It Fails Structurally	What the AI Actually Received
Bracketed stage directions	Typed [Callum is hiding fear beneath irritation] into the prompt	The AI treats brackets as metadata or converts them to narration; it does not build the feeling into rhythm, word choice, or pauses	A description of a feeling; no instructions for how the feeling reshapes speech
Revision-pass emotion injection	Generated dialogue first, then rewrote lines to add emotional weight	Emotion in dialogue is structural; it lives in the architecture of the sentence at generation time	Flat sentences with painted-on sentiment added afterward
Mood keywords	Typed "angry," "flirty," or "vulnerable" into the prompt	Keywords specify an emotional color but not how that color deepens, shifts, cracks, or bleeds through as the scene progresses	A one-note performance with no progression, no mask slip, no dynamic shift

The Core Principle: Stop describing feelings. Start constraining how feelings reshape speech. The Three-Tier Emotional Filter encodes emotion as structural instructions the AI executes during generation, producing subtext, tension, and progression in the output itself.